

MPI Project: Game of life

Tanuj Gupta

1. PCAM process

Summary:

I have implemented the Giga version(level 6) that decomposes the matrix in both directions.

1.1 Partitioning

- Partitioning involves decomposing computation and data into small tasks.
- To implement the Giga version and achieve maximum parallelism, I have decomposed the $N \times N$ grid in both directions. In other words, I have decomposed $N \times N$ grid into 2D subgrids
- By partitioning the $N \times N$ grid into 2D subgrids, we can assign a separate task to all these subgrids that can concurrently and independently solve the game of life.

1.2 Communication

We have decided to decompose the $N \times N$ grid into 2D subgrids. To be able to update the nodes on the edges, these sub grids need information from other nodes.

So every of these 2D subgrids would communicate to and from 8 neighbours:

A) Left and right Columns communication with 2 neighbours on the right and left

Edge cases: Ring communication

The subgrids on the left edge will communicate their left columns with the last subgrid in the same row. The subgrids on the right edge will communicate their right columns with the subgrids on the left edge will communicate their left

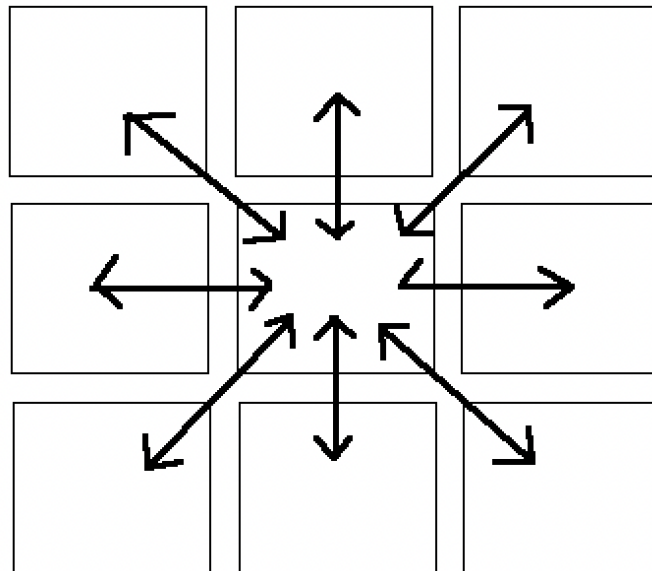
columns with

B) Top and bottom Rows communication with 2 neighbours on the bottom and Top

Edge cases: Ring communication

The subgrids on the top edge will communicate their top row with the last subgrid in the same column. The subgrids on the bottom edge will communicate their bottom row with the first subgrid in the same column

C) 4 corners communication with 4 neighbours on the diagonals.



Code Implementation for communication:

- Function `calculate_neighbour_nodes` in `utility_func.f90` calculates all the 8 neighbour subgrids for a given subgrid
- Function `communicate_columns_to_from_neighbours` in `utility_func.f90` communicates column information to and fro from 2 neighbours up and down
- Function `communicate_rows_to_from_neighbours` in `utility_func.f90` communicates row information to and fro from 2 neighbours left and right
- Function `communicate_diagnols_to_from_neighbours` in `utility_func.f90` communicates corner information to and fro from 4 neighbour on the diagonals

Other Communications:

1. Scatter subgrids to all nodes from nodes 0

Initially, the game data will be built on Node with rank 0. This node will have to scatter corresponding subgrids data to all the other nodes.

Code Implementation:

Functions `send_initial_grid_to_nodes` and `recieve_initial_grid_to_nodes` in `utility_func.f90` carry this communication.

2. Gather and aggregate updates on node 0 from all nodes

After update of the game locally on the 2D subgrid, all subgrids will send the update back to with rank 0.

Code Implementation:

Functions `send_updated_grid_to_pid0` and `recieve_updated_grid_to_pid0` in `utility_func.f90` carry this communication

1.3 Agglomeration

- We are decomposing the $N \times N$ grid into 2D subgrids. Theoretically as per partitioning, these subgrids can be as small as 1×1 size. Every subgrid will communicate with 8 of its neighbours. If we divide $N \times N$ grid into M such subgrids, total communication will be $8 \times M$.
- To reduce the amount of communication, we can group these smaller subgrids into larger ones. Larger size subgrids reduces communication costs and allows for tasks to be executed without interruption

1.4 Mapping

To perfectly balance the tasks across p processes, we need to assign equal sized grids to all processors. In an ideal scenario, every processor will get 2D subgrid of size

$\frac{(N \times N)}{P} = \frac{N}{\sqrt{P}} \times \frac{N}{\sqrt{P}}$. That is, we can assign each processor a square subgrid with dimension $\frac{N}{\sqrt{p}}$. But this is possible only if:

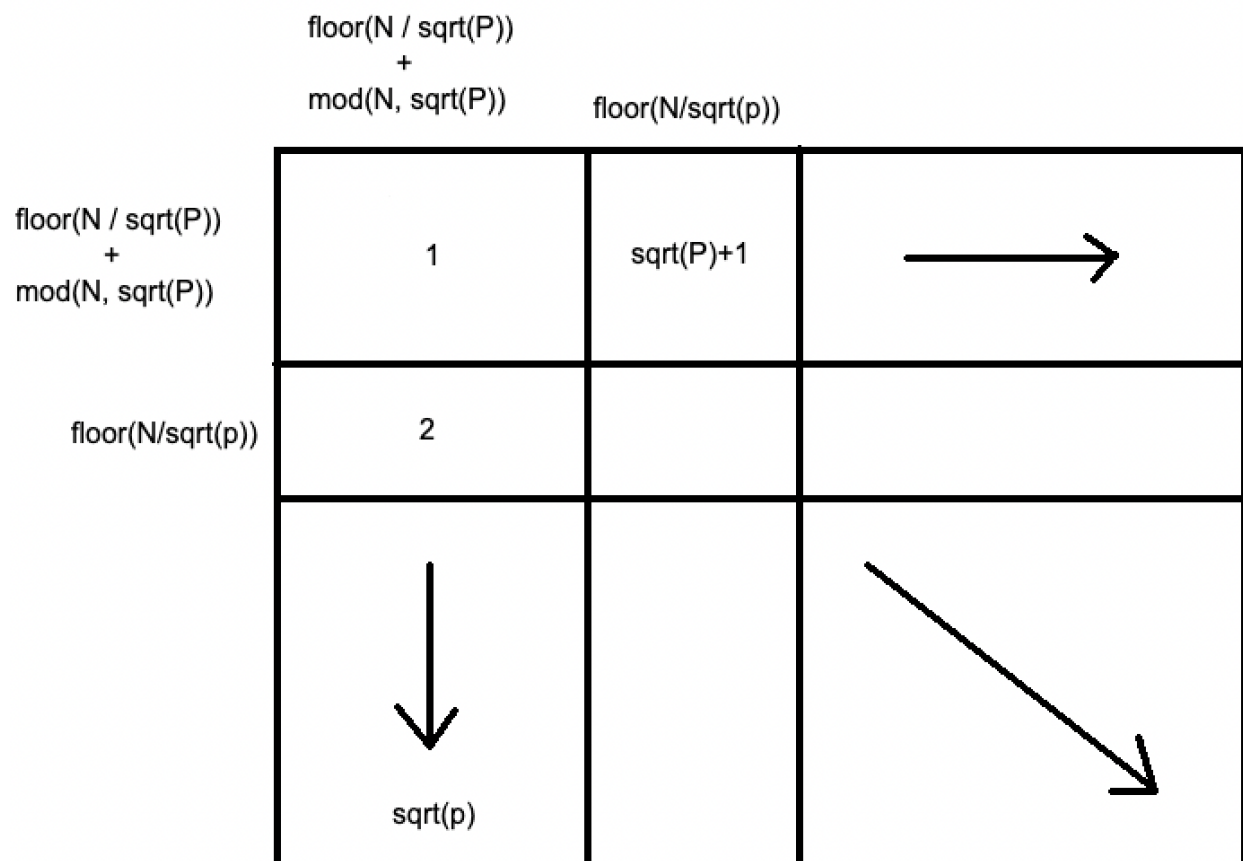
1. P is a perfect square.

To handle a more general scenario, we in our code factorize P into X and Y such that $P=X \times Y$. To make the division the most load balanced, we take X,Y such that $|X - Y|$ is minimum.

2. \sqrt{P} is divisible by N

To handle a more general scenario, where \sqrt{P} is not divisible by N , we assign the excess to the processors in first row and first column. In other words:

- The very first processor will be assigned a 2D subgrid of size $(\text{floor}(\frac{N}{\sqrt{P}}) + \text{mod}(N, P)) \times (\text{floor}(\frac{N}{\sqrt{P}}) + \text{mod}(N, P))$
- Processors in the first row will be assigned a 2D subgrid of size $(\text{floor}(\frac{N}{\sqrt{P}}) + \text{mod}(N, P)) \times (\frac{N}{\sqrt{P}})$
- Processors in the first column will be assigned a 2D subgrid of size $(\text{floor}(\frac{N}{\sqrt{P}})) \times (\text{floor}(\frac{N}{\sqrt{P}}) + \text{mod}(N, P))$
- All the other processors will be assigned a 2D subgrid of size $\text{floor}(\frac{N}{\sqrt{P}}) \times \text{floor}(\frac{N}{\sqrt{P}})$



Code Implementation for mapping:

- Function `calculate_local_grid_position` in `utility_func.f90` calculates relative position of processor on the grid in terms of row/col. Like processor 1 will have position 1,1. Processor P will have position \sqrt{P}, \sqrt{P}
- Function `calculate_local_grid_size` in `utility_func.f90` calculates the size of the 2D subgrid assigned to processor number n
- Function `calculate_local_grid_start_position` in `utility_func.f90` calculates the indexes of first cell of the 2D subgrid assigned to processor number n .

2. How to run the code

On grape, do the following to compile and run my code:

STEP 1: Make sure all the required libraries are loaded using commands:

1. `module load rocks-openmpi`
2. `module load gcc-6.2.0`

STEP 2: To change the

- A) Grid size: Edit `game_of_life_mpi_mega_version.f90` and change variable `grid_size`
- B) The number of processors: Edit `game_of_life_mpi_mega_version.f90` and change variable `num_procs`
- C) The number of iterations: Edit `game_of_life_mpi_mega_version.f90` and change variable `num_of_iterations`

By default, values of these variables are:

`grid_size=20`

`num_procs=4`

`num_of_iterations=80`

STEP 3: To compile my code, run the following commands in order

```
mpif90 -o utility_func.o -c utility_func.f90
```

```
mpif90 -o game_of_life_mpi_giga_version.o utility_func.o game_of_life_mpi_giga_version.f90
```

STEP 4: To run my code with n processors, run the following command

```
mpirun -np n game_of_life_mpi_giga_version.o
```

3 Code output for 20×20 Grid

[illegible]

Grid After iteration:

20

A 20x20 grid with the number 1 in the 10th column and 10th row, and the number 1 in the 11th column and 10th row.

Grid After iteration:

40

A 20x20 grid of dots. The number '1' is placed in the center, specifically at the intersection of the 10th column and the 10th row, representing the identity matrix I_2 .

Grid After iteration:

60

Grid After iteration:

80

4. Performance Model and Testing

Assume

- **Dimensions** = $N_x \times N_y$, where $N_x = N_y = N$ for simplicity
- **Number of processes** = P
- 2-D domain decomposition

Given the dimensions:

$$\text{Total grid points} = N \times N = N^2 \quad - (1)$$

Computation time

Assuming No replicated computation:

$$T_{comp} = T_c \times N^2 \quad - (2)$$

where t_c = average computation time per grid point (slightly different at edges from interior etc)

Communication time

Given that its a 2-D domain decomposition, we will partition the grid points among P processors in 2 directions, namely, X and Y.

Assuming that the partitioning is equal in both the directions, **the dimension of task per processor** would then be:

$$\text{X direction: } \frac{N}{\sqrt{P}}, \quad \text{Y direction: } \frac{N}{\sqrt{P}}, \quad \text{Z direction: } N_z$$

$$\text{Total grid points per processor} = \frac{N}{\sqrt{P}} \times \frac{N}{\sqrt{P}} \quad - (3)$$

Given that its a 2d decomposition, each task communicates with 8 neighbours. With 4 of its neighbours, the size of the wall shared is $\frac{N}{\sqrt{P}}$. With 4 diagonal neighbours, the size of the wall shared is 1.

Therefore:

$$\begin{aligned} T_{comm} &= 4P * (T_s + T_w * \frac{N}{\sqrt{P}}) + 4P * (T_s + T_w) \\ \Rightarrow T_{comm} &= 4P * (T_s + T_w * (\frac{N}{\sqrt{P}} + 1)) \quad - (4) \end{aligned}$$

4.1. Performance Model: 2D Game of life Total time

Assuming that P divides N exactly, then assume load-balanced and no idle time:

$$T_{2D_gol} = (T_{comp} + T_{comm})/P$$

$$\Rightarrow T_{2D_gol} = \frac{T_c N^2}{P} + 4P * \left(\frac{T_s + T_w * (\frac{N}{\sqrt{P}} + 1)}{P} \right)$$

$$\Rightarrow T_{2D_gol} = \frac{T_c N^2}{P} + 4(T_s + T_w * (\frac{N}{\sqrt{P}} + 1))$$

$$\Rightarrow T_{2D_gol} = \frac{T_c N^2}{P} + 4T_s + 4T_w * (\frac{N}{\sqrt{P}} + 1) \quad - (5)$$

4.2. Values for T_s , T_w , and T_c

Values for T_s , T_w

From my HW6, values obtained for T_s and T_w were:

$$T_s = 0.007997999 \text{ sec}$$

$$T_w = 0.0005823312 \text{ sec}$$

Values for T_c

To calculate value for constant T_c , I used ones.f90 code used for HW2.

Basically

$$T_{comp} = T_c \times N^2$$

$$\Rightarrow T_c = \frac{T_{comp}}{N^2} \quad -(6)$$

Using eq(6), we can estimate T_c . Idea is to compute (T_{comp}/N^2) for different values of N to estimate T_c . Then take average of these values to get accurate estimation of T_c

Data Point	Value of N	Value of T_{comp}	Value of $T_c = T_{comp}/N^2$
1	500	0.0339939	$1.359759 * 10^{-7}$
2	1000	0.132979	$1.32979 * 10^{-7}$

3	3000	1.1698219	$1.299802 * 10^{-7}$
4	6000	4.6652909	$1.29591416 * 10^{-7}$
5	9000	10.4964039	$1.295852345679 * 10^{-7}$
6	15000	29.159567	$1.29598075 * 10^{-7}$
Average -> $\frac{1}{6} \left(\sum_{n=1}^6 T_c \right) / 6$			$1.31284970928 * 10^{-7}$

$$\Rightarrow T_c = 1.31284970928 * 10^{-7}$$

$$\Rightarrow T_s = 7.997999 * 10^{-3}$$

$$\Rightarrow T_w = 5.823312 * 10^{-4}$$

4.3. Testing

D.1 Time estimation using Performance model given by equation 5

For data points:

N= 20, 100, 500, 1000, 3000, 6000

As per equation 5

$$\Rightarrow T_{2D_gol} = \frac{T_c N^2}{P} + 4T_s + 4T_w * \left(\frac{N}{\sqrt{P}} + 1 \right)$$

Using

$$P = 4,$$

$$T_c = 1.31284970928 * 10^{-7},$$

$$T_s = 7.997999 * 10^{-3},$$

$$T_w = 5.823312 * 10^{-5}$$

Data Point	Value of N	Value of T_{2D_gol}
1	20	0.0576276971187326126

2	100	0.15111577458446845
3	500	0.62485783733427525
4	1000	1.2318049520254135
5	3000	3.8236998692154884
6	6000	8.2038607969880104

D.2 Actual Running Time using the code

```
[~bash-4.1$ mpirun -np 5 game_of_life_mpi_giga_version_timings.o
DAT Registry: sysconfdi, bad filename - /etc/rdma/compat-dapl/dat.conf, retry default at /etc/dat.conf
DAT Registry: sysconfdi, bad filename - /etc/rdma/compat-dapl/dat.conf, retry default at /etc/dat.conf
DAT Registry: sysconfdi, bad filename - /etc/rdma/compat-dapl/dat.conf, retry default at /etc/dat.conf
DAT Registry: sysconfdi, bad filename - /etc/rdma/compat-dapl/dat.conf, retry default at /etc/dat.conf
DAT Registry: sysconfdi, bad filename - /etc/rdma/compat-dapl/dat.conf, retry default at /etc/dat.conf
For N:      20
For P:      4
Total time taken:  6.9979999999999728E-003
-----
For N:      100
For P:      4
Total time taken:  2.5996999999999993E-002
-----
For N:      500
For P:      4
Total time taken:  0.30895399999999995
-----
For N:      1000
For P:      4
Total time taken:  1.2948029999999999
-----
For N:      3000
For P:      4
Total time taken:  11.454257999999999
-----
For N:      6000
For P:      4
Total time taken:  41.629671999999999
-----
```

Data Point	Value of N	Value of T_{2D_gol}
1	20	0.00699799999999999728
2	100	0.0259969999999999993
3	500	0.30895399999999995
4	1000	1.2948029999999999
5	3000	11.454257999999999

6	6000	41.629671999999999
---	------	--------------------

D.3 Comparison

Data Point	Value of N	Value of Time(Estimated by Model)	Value of Time(Actual)
1	20	0.0576276971187326126	0.0069979999999999728
2	100	0.15111577458446845	0.025996999999999993
3	500	0.62485783733427525	0.30895399999999995
4	1000	1.2318049520254135	1.2948029999999999
5	3000	3.8236998692154884	11.454257999999999
6	6000	8.2038607969880104	41.629671999999999

