

HW6

1. TASK 1: PING PONG LATENCY

Summary:

I have used 1000 iterations of ping-pong loop to get accurate timings. Finally ran the code for 16 data points with data following message_length(L) in bytes: 0, 100, 1000, 10000, 100000, 500000, 1000000, 2000000, 3000000, 4000000, 5000000, 6000000, 7000000, 8000000, 9000000, 10000000. The code calculates total timing Tmsg for all the above mentioned data points.

To calculate T_s , I simply used the value of Tmsg for message_length(L) = 0.

To calculate T_w , I am using mean of T_w value for all the 16 data points.

$T_s = 0.007997999 \text{ sec}$

$T_w = 0.00005823312 \text{ sec}$

Loop Iterations used:


1000

Data Points:

Data Points	X Axis -> Message Length(L)	Y Axis -> T_{msg}
1	0	0.007997999
2	100	0.013313002
3	1000	0.048486501
4	10000	0.363008
5	100000	4.0599218
6	500000	34.8613816
7	1000000	51.0772118

8	2000000	132.3538463
9	3000000	179.9873296
10	4000000	297.8985994
11	5000000	351.6244534
12	6000000	427.94704218
13	7000000	441.8691686
14	8000000	481.82289031
15	9000000	539.689148131
16	10000000	574.63355542

Graph:



Calculation for T_s :

- For $L=0$, $T_{msg} = T_s$
- From above data points, T_{msg} for $L=0$ is 0.007997999 sec
- Therefore, $T_s = 0.007997999$ sec

Calculation for T_w :

- For any data point (L , T_{msg})
- $$T_{msg} = T_s + T_w * L$$
- $$\Rightarrow T_w = (T_{msg} - T_s) / L$$

Data Points	X Axis -> Message Length(L)	Y Axis -> T_{msg}	$(T_{msg} - T_s)/L$
1	0	0.007997999	
2	100	0.013313002	0.00005315002000000007 10
3	1000	0.048486501	0.00004048850099999982 79
4	10000	0.363008	0.00003550099999999995 97
5	100000	4.0599218	0.00004051923799999984 07
6	500000	34.8613816	0.00006970676720000000 78
7	1000000	51.0772118	0.00005106921379999996 89
8	2000000	132.3538463	0.00006617292415000002 06
9	3000000	179.9873296	0.00005999311053333331 51
10	4000000	297.8985994	0.00007447265034999998 89
11	5000000	351.6244534	0.00007032329
12	6000000	427.94704218	0.00007132317
13	7000000	441.8691686	0.00006312302437142862 4
14	8000000	481.82289031	0.00006022686
15	9000000	539.689148131	0.00005996457223677777 879
16	10000000	574.63355542	0.00005746255574200000 90
SUM -> $\sum_{n=1}^{15} (T_{msg} - T_s)/L$			0.00087349689

- Now, we will take mean of all the values for all data points m to get most accurate value for T_w , i.e,

$$T_w = \left(\sum_{n=1}^m (T_{msg} - T_s) / L \right) \div m$$

$$T_w = \left(\sum_{n=1}^{15} (T_{msg} - T_s) / L \right) \div m$$

$$T_w = \left(\sum_{n=1}^{15} (T_{msg} - T_s) / L \right) \div 15$$

$$T_w = \left(\sum_{n=1}^m (T_{msg} - T_s) / L \right) \div 15$$

$$T_w = 0.00087349689 \div 15$$

$$T_w = 0.00005823312 \text{ sec}$$

2. TASK 2: FINITE-DIFFERENCE FOR 2D DATA DECOMPOSITION

Assume

- **Dimensions** = $N_x \times N_y \times N_z$, where $N_x = N_y = N$ for simplicity
- 9-point stencil
- **Number of processes** = P
- 2-D domain decomposition

Given the dimensions:

$$\text{Total grid points} = N \times N \times N_z = N^2 N_z \quad - (1)$$

Calculations:

A. Computation time

Assuming No replicated computation:

$$T_{comp} = t_c \times N^2 \times N_z \quad - (2)$$

where t_c = average computation time per grid point (slightly different at edges from interior etc)

B. Communication time

Given that its a 2-D domain decomposition, we will partition the grid points among P processors in 2 directions, namely, X and Y.

Assuming that the partitioning is equal in both the directions, **the dimension of task per processor** would then be:

$$\text{X direction: } \frac{N}{\sqrt{P}}, \quad \text{Y direction: } \frac{N}{\sqrt{P}}, \quad \text{Z direction: } N_Z$$

$$\text{Total grid points per processor} = \frac{N}{\sqrt{P}} \times \frac{N}{\sqrt{P}} \times N_Z \quad - (3)$$

Given that its a 9 point stencil and 2d decomposition, each task exchanges 2 planes with each of 4 neighbours.

Size of walls with these 4 neighbours is:

$$\frac{N}{\sqrt{P}} \times N_Z$$

Therefore:

$$T_{comm} = 4P * (T_S + 2 * T_W * \frac{N}{\sqrt{P}} * N_Z) \quad - (4)$$

C. 2D Finite difference time

Assuming that P divides N exactly, then assume load-balanced and no idle time:

$$T_{2D_finite_diff} = (T_{comp} + T_{comm})/P$$

$$\Rightarrow T_{2D_finite_diff} = \frac{t_c N^2 N_Z}{P} + 4P * \left(\frac{T_S + 2 * T_W * \frac{N}{\sqrt{P}} * N_Z}{P} \right)$$

$$\Rightarrow T_{2D_finite_diff} = \frac{t_c N^2 N_Z}{P} + 4 * (T_S + 2 * T_W * \frac{N}{\sqrt{P}} * N_Z)$$

$$\Rightarrow T_{2D_finite_diff} = \frac{t_c N^2 N_Z}{P} + 4T_S + 8T_W * \frac{N}{\sqrt{P}} * N_Z \quad - (5)$$

D. 2D Finite difference Efficiency

We know, relative efficiency is given by:

$$E_{relative} = \frac{T_1}{PT_p} \quad - (6)$$

Where, T_1 is time to run algorithm by a single processor, and T_p is time to run algorithm on P processors

For our algorithm,

$$T_1 = t_c N^2 N_Z \quad \text{And} \quad T_p = T_{2D_finite_diff} = \frac{t_c N^2 N_Z}{P} + 4T_S + 8T_W * \frac{N}{\sqrt{P}} * N_Z$$

$$\Rightarrow E = \frac{T_1}{PT_P} = \frac{t_c N^2 N_Z}{P * (\frac{t_c \times N^2 \times N_Z}{P} + 4T_S + 8T_W * \frac{N}{\sqrt{P}} * N_Z)}$$

$$\Rightarrow E = \frac{t_c N^2 N_Z}{t_c N^2 N_Z + 4PT_S + 8\sqrt{P} T_W N N_Z} \quad - (7)$$

E. 2D Finite difference Iso-Efficiency

For 2D decomposition, efficiency we calculated in equation 7 is:

$$\Rightarrow E = \frac{t_c N^2 N_Z}{t_c N^2 N_Z + 4PT_S + 8\sqrt{P} T_W N N_Z}$$

$$\Rightarrow t_c N^2 N_Z = E * (t_c N^2 N_Z + 4PT_S + 8\sqrt{P} T_W N N_Z)$$

Here, dominant terms are:

$$\sim N^2 N_Z \approx \sim N^2 N_Z \quad \sim \sqrt{P} N N_Z$$

For E to be constant,

$$N^2 N_Z = \sqrt{P} N N_Z$$

$$\Rightarrow N = \sqrt{P} \quad - (8)$$

By substituting $N = \sqrt{P}$, E is constant given by equation

$$t_c N_Z = E * (t_c N_Z + 4T_S + 8T_W N_Z) \quad - (9)$$

Total computation T_{comp} , given by equation 2 is $T_{comp} = t_c N^2 N_Z$

Therefore, to keep E constant: $T_{comp} \propto N^2$

$$\Rightarrow T_{comp} \propto P \text{ (Using Eq 8: } N = \sqrt{P} \text{)} \quad - (10)$$

Number of grid points, given by equation 1 is $Grid\ points = N^2 N_Z$

Therefore, to keep E constant: $Grid\ points \propto N^2$

$$\Rightarrow \text{Grid points} \propto P \text{ (Using Eq 8: } N = \sqrt{P} \text{)} \quad - (11)$$

From equations 10 and 11, we can conclude:

the isoefficiency of this algorithm is $O(P)$

$$\Rightarrow \text{Isoefficiency} = O(P) \quad - (12)$$

F. Analysis

Comparative Analysis of 2D_Decomposition and 1D_Decomposition

- 2D_Decomposition with *Isoefficiency* of $O(P)$ is way more scalable than 1D_Decomposition which has *Isoefficiency* of $O(P^2)$.
- This is because the amount of computation needed to keep E constant in 2D_Decomposition must increase linearly with the number of processors to keep the efficiency constant. This is way better than 1D_Decomposition, where the amount of computation needed to keep E constant must increase as the square of the number of processors.

Reasoning:

- The number of neighbors in 2D_Decomposition taking part in the communication is 4 per processor, 2 more than the number of neighbors in 1D_Decomposition per processor.
- Although the amount of communication increases from $2P$ to $4P$, the number of messages/grid_points shared per communication reduces.
- This is because the wall area is lesser for 2D_Decomposition ($\frac{N}{\sqrt{P}} * N_z$) in comparison to 1D_Decomposition's wall area ($N * N_z$).

Conclusion:

- Decomposition that increases the number of walls, will end in better efficiency. This is because, although the amount of communication may increase, the number of messages/grid_points shared per communication reduces.